

Tweet id reconstruction challenge

In the following short report, I will show my analysis of the random tweet ids.

A tweet id consists of 44-bit binary number that is composed of 4 parts: Time component (22 bits), Data Center Id (5 bit), Server Id (5 bit), Sequence Id (12 bit).

I will use a combination of the Data Center Id and Server Id as the machine id.

Initial pre-scan results

After examining the random ids file, I found that there are total of 81 sequence ids, and 38 total machines ids used in the random set. And the top four machines used are responsible to process 20% of the tweets.

furthermore, after examining the sequence ids I found that 7 ids were enough to get over 99% of the tweets and that using sequence id 0 will yield 61% of them.

Machine id	%	Total %	Sequence id	%	Total %
363	0.05437	0.05437	0	0.61148	0.61148
375	0.05402	0.10839	1	0.2363	0.84778
382	0.05374	0.16213	2	0.07281	0.9206
365	0.0507	0.21283	5	0.02683	0.94743
366	0.04918	0.26201	3	0.01914	0.96657
381	0.04839	0.3104	6	0.0165	0.98307
336	0.04496	0.35536	7	0.0066	0.98967
361	0.04488	0.40024	4	0.00453	0.9942
350	0.04486	0.4451	8	0.00209	0.9963
335	0.04485	0.48995	10	0.00116	0.99746
333	0.04468	0.53463	11	0.00086	0.99832
373	0.03506	0.5697	9	0.0006	0.99893
377	0.03347	0.60317	12	0.00043	0.99936
332	0.03304	0.6362	13	0.00017	0.99953
352	0.03097	0.66718	15	8e-05	0.99961
340	0.02545	0.69263	14	8e-05	0.99969
327	0.02524	0.71787	16	7e-05	0.99975
339	0.02519	0.74306	17	4e-05	0.99979
376	0.02417	0.76723	18	2e-05	0.99982
364	0.02262	0.78985	19	2e-05	0.99984
347	0.02003	0.80988	20	2e-05	0.99986
326	0.0196	0.82948	32	1e-05	0.99987
324	0.01884	0.84832	21	1e-05	0.99988
348	0.01878	0.8671	22	1e-05	0.9999
325	0.01843	0.88553	33	1e-05	0.99991
334	0.0171	0.90263	24	1e-05	0.99991
378	0.01566	0.91829	23	1e-05	0.99992
323	0.01353	0.93182	25	1e-05	0.99993
379	0.01223	0.94405	27	1e-05	0.99994
346	0.01172	0.95578	26	1e-05	0.99994
362	0.01092	0.96669	28	1e-05	0.99995
372	0.01079	0.97749			
342	0.01003	0.98752			
321	0.00689	0.99441			
349	0.00515	0.99956			
322	0.00021	0.99977			
328	0.00013	0.9999			
320	0.0001	1.0			
total machines: 38					

Figure 1: Machine ids

Figure 2: Sequence ids

Naive method analysis

Combining these two facts I can deduce that while there are 4096 different sequence ids per machine, we can only check 7 of the values to get more than 99% of all the messages, and since the machine ids are 10 bit long we get from 1024 different machine ids to 34 to get over 99% of the tweets.

This is an improvement by a factor of $585 * 30 = 17,550$.

Now if we want only 10% of the full FireHose, we will perform the following calculation:

We have 3600 API calls per hour → 360,000 tweets per hour → 90,000 tweets per 15 minutes

Given that we have $38 \text{ machines} * 7 \text{ sequence numbers} = 266 \text{ scans needed per millisecond}$

We have access to 90,000 tweets every 15 minutes → $\frac{90,000}{266} * 10 = 3383.4 \text{ ms} (3.38 \text{ seconds})$ of data every 150 minutes

So we can get $\frac{24 \text{ hour} * 60 \text{ minutes} * 3.38 \text{ s data}}{150} = 32.48 \text{ seconds of data per day of constant API calls}$

Since we want 10% of the data we will need:

$$\frac{24 * 60 * 6}{32.48} = 266 \text{ keys} \rightarrow 0.266 \text{ API calls per millisecond} \rightarrow$$

$90,000 * 266 = 23,940,000 \text{ API calls}$ for a single account just to get that 10%.

Improved method

If we will only scan the top 4 machines to get 21% of the tweets, and sequence id of 0 we will get ~12% of the tweets. By following the same calculation, we will get that every single account will yield 225 seconds of data every 150 minutes. Following this we can get for every day of constant calls 2160 seconds of the tweets. This gives us the following:

$$\frac{86400 \text{ seconds per day}}{2160 \text{ seconds of data per day}} = 40 \text{ API keys} \rightarrow 0.04 \text{ API calls per millisecond}$$

This is an improvement by a factor of 60 from the naïve approach of scanning 7 sequence ids per machine, and will require 6 time less keys.

After examining the 1k, 10k, 100k files with sequence 0 we get the following:

	1k	10k	100k	average
363	29.925	40.149	45.971	38.68167
375	31.164	39.538	45.696	38.79933
382	31.150	40.058	45.202	38.80333
365	29.728	41.110	45.951	38.92967

So, while using the same API calls, we will get on average 38.8035% of the busiest tweets, from the FireHose on top of the 12% of the entire FireHose.

Future Work

This short analysis could be expended further in a few different ways.

First, it would be possible to find an optimal value of the number of ids that needs to be scanned as a function of % of total tweets received, as a cross product of the random sample along with the top 1k,10k,100k busiest minutes.

Secondly, given a longer period and access to several developer accounts it would be possible to check this analysis and verify that the numbers correspond and yield good results while matched with the suggested method.

Lastly, train a model using the top milliseconds data and use this model to predict when peaks will occur and capture them without the need to constantly perform API calls to twitter.